

PDCA or CIPDA?

A guide for professionals deciding whether and how to use the CIPDA framework

Read this before the Quick-Start Guide. It will tell you whether CIPDA is right for your situation — and at what depth.

1. Start here

PDCA has not stopped working. The question is whether it is working well enough when AI is part of the analytical process. This guide helps you answer that for your specific situation — before you invest time in learning a new framework.

Three outcomes are possible after reading this:

- PDCA is sufficient for your AI use — carry on.
- CIPDA would help but a light version is enough — start there.
- Your situation calls for the full CIPDA framework — read the Quick-Start Guide.

2. What PDCA does well — and where it strains

What PDCA does well — and why it still matters

PDCA (Plan–Do–Check–Act) has been the backbone of quality and improvement management for decades. It works because it imposes a discipline on human thinking: slow down, plan before acting, check what happened, learn from it.

It remains highly effective when:

- Human practitioners share the same context, vocabulary, and institutional knowledge
- Tasks are well-defined and outputs are verifiable by a competent reviewer
- The analytical work is routine and the decision logic is stable
- Speed and pragmatism matter more than full traceability

If your AI use fits this description, PDCA is sufficient. You do not need CIPDA.

Where PDCA starts to strain when AI is involved

PDCA was designed for human practitioners who share situational awareness. When you plan a task with a colleague, you both carry implicit context — what the organisation values, what the constraints are, what 'good' looks like. That shared understanding is invisible but essential.

AI does not carry that context. It responds probabilistically to what it is given. If the framing is incomplete or ambiguous, the AI fills the gaps with its own assumptions — and those assumptions may differ each time the task is run.

The practical consequence: research testing 375 identical analytical tasks found that unstructured prompting produced consistent results only 74% of the time. PDCA-structured prompting gave 75%. That 1% gap tells you PDCA alone is not the answer — structuring the output stage without structuring the reasoning stage does not suppress the underlying variance.

The strain shows up as:

- Different answers to the same question across sessions
- Outputs that look reasonable but rest on assumptions you never made
- Reasoning you cannot inspect or defend if challenged
- Constraint violations that only become visible after the output is used

None of these are failures of the AI model. They are failures of reasoning architecture.

3. Three levels of engagement

CIPDA is not all-or-nothing. You can engage with it at the level your task requires.

	Level 1 — PDCA as usual	Level 2 — CIPDA light	Level 3 — CIPDA full
What it is	Standard PDCA applied around AI outputs. AI assists with execution; humans plan and check as normal.	CIPDA Context and Intent stages added to front-load the AI's framing. No formal Directive. ROMER optional.	Full Primer with Directive Statement, all five CIPDA stages built interactively, and ROMER trace at Assure.
Use it when	• Routine, well-defined tasks • Output is easily verified • Low accountability exposure • Team shares context	• AI is doing meaningful analytical work • Outputs will inform a decision • Some consistency matters but stakes are moderate	• Competing or hard constraints exist • Output will be shared, signed off, or relied on • Regulatory, audit, or compliance exposure • Same result needed every time • Reasoning must be inspectable
What you add	Nothing new. Use your existing process.	A short Context and Intent statement before your prompt. Takes 5–10 minutes.	A Directive Statement (written offline) plus all five Primer stages built interactively with the AI. Plan for 20–40 minutes for a new task type; much less once you have a template.
The trade-off	Fast. No overhead. But inconsistent on complex or constraint-sensitive tasks.	Moderate improvement in consistency. Reasoning still partially implicit.	Near-complete consistency (97% in controlled testing). Full reasoning trail. Higher upfront effort, reusable for repeat tasks.

4. Read the signals in your task

Work through these questions for the specific task in front of you. The answers will point you to a level.

Quick read — where does your task sit?	
If you ask yourself...	The signal is...
Could I verify this output is right without running it again?	Yes → Level 1 is fine
Would it matter if the AI reached a different conclusion next time?	Yes → consider Level 2
Will someone other than me rely on or act on this output?	Yes → consider Level 2 or 3
Are there hard constraints the AI must not breach?	Yes → Level 3

Do I need to show my reasoning if this decision is challenged?	Yes → Level 3
Is personal, commercial, or regulated data involved?	Yes → Level 3 with Directive
Would the consequences of a wrong answer be serious or hard to reverse?	Yes → Level 3 with Directive

5. On using it superficially

A note on superficial use

You may decide to use CIPDA lightly — just the Context and Intent stages, without a Directive or formal ROMER trace. That is a reasonable choice for lower-stakes work and it will improve your results over unstructured prompting.

Be clear with yourself about what you are trading away. Superficial use improves consistency but does not provide the governance assurance that a full Primer and ROMER trace give you. If the task later turns out to carry more weight than you expected — it gets shared, challenged, or acted on — a partial Primer will not protect you.

The right question is not ‘how little can I do?’ but ‘what does this task actually require?’ Answer that honestly and the level chooses itself.

6. What to do next

What to do next	
You chose Level 1	No action needed beyond your existing PDCA practice. If a task later grows in complexity or accountability, revisit this guide.
You chose Level 2	Before your next AI-assisted task, write two sentences: what the task is and what the non-negotiables are. Paste them at the top of your prompt before anything else. That is a Level 2 Primer.
You chose Level 3	Open the CIPDA + ROMER Quick-Start Guide. Start at Step 0. Build your Directive Statement before you open the AI session.

“PDCA governs AI outputs from the outside; it does not govern the reasoning process from within. Where AI systems participate directly in analytical work, this outside-in posture leaves interpretive variance unaddressed at its source.”

— Kenneth Tombs, IJQRM (2026)

Document information

Version: 1.0 **Date:** June 2026

This document should be read alongside the CIPDA + ROMER Quick-Start Guide for Professionals.

Founding article:

Tombs, K. (2026). From PDSA to PDCA to CIPDA: structured improvement in AI-supported management reasoning. *International Journal of Quality & Reliability Management*. DOI: 10.1108/IJQRM-01-2026-0029

Copyright

© 2026 Kenneth Tombs. This guide, including the CIPDA framework, Directive Statement structure, and ROMER mechanism as presented herein, is the intellectual property of Kenneth Tombs. Use is permitted for academic

research, education, and commercial application provided appropriate attribution is maintained. This statement does not transfer ownership of the underlying architecture, know-how, or related derivative systems.

© **2026 Emerald Publishing Limited**. The founding article — Tombs, K. (2026), 'From PDSA to PDCA to CIPDA: structured improvement in AI-supported management reasoning', *International Journal of Quality & Reliability Management* — is published by Emerald Publishing Limited. All rights in the published article are reserved. Reproduction of the article in whole or in part requires permission from Emerald Publishing Limited. DOI: 10.1108/IJQRM-01-2026-0029.